

Predictive Modeling for Amazon Prime Subscription Plans

May 26, 2024

1 Predictive Modeling for Amazon Prime Subscription Plans: Insights from User Data

1.1 Introduction

This project focuses on analyzing Amazon subscriber data to build predictive models that determine factors influencing subscription plan choices (Annual vs. Monthly). By leveraging machine learning techniques, the goal is to provide insights that can help Amazon optimize its subscription offerings, improve customer retention, and enhance user satisfaction.

1.2 Project Workflow

1. **Data Preprocessing:** Cleaning and preparing the dataset for analysis.
2. **Exploratory Data Analysis (EDA):** Understanding the data through visualization and summary statistics.
3. **Model Development:** Building and evaluating different machine learning models.
4. **Model Selection:** Choosing the best-performing model based on evaluation metrics.
5. **Conclusion and Recommendations:** Summarizing findings and providing actionable insights.

The dataset includes various features related to Amazon subscribers, such as user demographics, usage patterns, and engagement metrics. Below is a summary of the key features:

Field	Description
User ID	Numeric ID for the user
Name	User's name
Email Address	User's email address
Username	Username
Date of Birth	Date of Birth of the user
Gender	User's gender
Location	User's Country
Membership Start Date	Start date of the Amazon Prime membership
Membership End Date	End date of the membership
Payment Information	Payment method used (Visa, Mastercard, Amex)
Renewal Status	Setting of renewal on subscription (Manual vs Automatic)
Usage Frequency	Frequency of platform usage (Occasional, Regular, Frequent)
Purchase History	Most frequently purchased items
Favorite Genres	Favorite genre of content

Field	Description
Devices Used	Device used to access the platform
Engagement Metrics	Engagement level (Low, Medium, High)
Feedback/Ratings	Average feedback/ratings given by the user
Customer Support Interactions	Number of interactions with customer support
Subscription Plan	Type of subscription plan (Annual vs Monthly)

1.3 Imports

```
[4]: # Import Libraries

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from datetime import datetime

from sklearn.model_selection import train_test_split, PredefinedSplit, GridSearchCV
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import make_scorer, accuracy_score, precision_score, recall_score, classification_report, confusion_matrix, f1_score

from xgboost import XGBClassifier, plot_importance

import warnings

# Ignore specific FutureWarning related to seaborn
warnings.simplefilter(action='ignore', category=FutureWarning)
```

2 Data Preprocessing

```
[6]: # Load the dataset
data = pd.read_csv("C:/Users/luke3/Documents/GitHub/
    Predictive-Modeling-for-Amazon-Prime-Subscription-Plans-Insights-from-User-Data/
    data/amazon_prime_users.csv")

[7]: # Display the first 5 rows of the data set
data.head()
```

```
[7]:
```

	User ID	Name	Email Address \
0	1	Ronald Murphy	williamholland@example.com
1	2	Scott Allen	scott22@example.org
2	3	Jonathan Parrish	brooke16@example.org
3	4	Megan Williams	elizabeth31@example.net
4	5	Kathryn Brown	pattersonalexandra@example.org

	Username	Date of Birth	Gender	Location \
0	williamholland	1953-06-03	Male	Rebeccachester
1	scott22	1978-07-08	Male	Mcpersonview
2	brooke16	1994-12-06	Female	Youngfort
3	elizabeth31	1964-12-22	Female	Feliciashire
4	pattersonalexandra	1961-06-04	Male	Port Deborah

	Membership Start Date	Membership End Date	Subscription Plan \
0	2024-01-15	2025-01-14	Annual
1	2024-01-07	2025-01-06	Monthly
2	2024-04-13	2025-04-13	Monthly
3	2024-01-24	2025-01-23	Monthly
4	2024-02-14	2025-02-13	Annual

	Payment Information	Renewal Status	Usage Frequency	Purchase History \
0	Mastercard	Manual	Regular	Electronics
1	Visa	Manual	Regular	Electronics
2	Mastercard	Manual	Regular	Books
3	Amex	Auto-renew	Regular	Electronics
4	Visa	Auto-renew	Frequent	Clothing

	Favorite Genres	Devices Used	Engagement Metrics	Feedback/Ratings \
0	Documentary	Smart TV	Medium	3.6
1	Horror	Smartphone	Medium	3.8
2	Comedy	Smart TV	Low	3.3
3	Documentary	Smart TV	High	3.3
4	Drama	Smart TV	Low	4.3

	Customer Support Interactions
0	3
1	7
2	8
3	7
4	1

```
[8]: # Check data types
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2500 entries, 0 to 2499
Data columns (total 19 columns):
```

#	Column	Non-Null Count	Dtype
0	User ID	2500 non-null	int64
1	Name	2500 non-null	object
2	Email Address	2500 non-null	object
3	Username	2500 non-null	object
4	Date of Birth	2500 non-null	object
5	Gender	2500 non-null	object
6	Location	2500 non-null	object
7	Membership Start Date	2500 non-null	object
8	Membership End Date	2500 non-null	object
9	Subscription Plan	2500 non-null	object
10	Payment Information	2500 non-null	object
11	Renewal Status	2500 non-null	object
12	Usage Frequency	2500 non-null	object
13	Purchase History	2500 non-null	object
14	Favorite Genres	2500 non-null	object
15	Devices Used	2500 non-null	object
16	Engagement Metrics	2500 non-null	object
17	Feedback/Ratings	2500 non-null	float64
18	Customer Support Interactions	2500 non-null	int64

dtypes: float64(1), int64(2), object(16)
memory usage: 371.2+ KB

```
[9]: # Convert columns to datetime
date_columns = ['Date of Birth', 'Membership Start Date', 'Membership End Date']
data[date_columns] = data[date_columns].apply(pd.to_datetime)
```

```
[10]: ## Check to see if any email addresses are duplicated
duplicate_email_rows = data[data.duplicated(subset=['Email Address'],
↳keep=False)]

duplicate_email_rows.head(10)

# Because this is a sample dataset and the Names are different for each
↳duplicated email address or username, I'll ignore this and leave them all in.
```

```
[10]:
```

	User ID	Name	Email Address	Username \
6	7	Benjamin Marshall	michaellewis@example.net	michaellewis
33	34	Whitney Underwood	ujones@example.com	ujones
84	85	Keith Baker	ujones@example.com	ujones
107	108	Grant Jensen	tyler29@example.com	tyler29
120	121	Lisa Washington	dbailey@example.net	dbailey
129	130	Jordan Jackson	ubrown@example.org	ubrown
147	148	Troy Smith	john96@example.net	john96
175	176	Alexandra James	twilliams@example.com	twilliams
401	402	Olivia Harper	twilliams@example.com	twilliams

443 444 Brent Key sburke@example.com sburke

	Date of Birth	Gender	Location	Membership Start Date	\
6	2003-02-09	Male	Carlsonchester	2024-04-08	
33	2002-07-23	Female	New Jenniferport	2024-03-30	
84	1972-07-11	Female	Langton	2024-02-03	
107	1944-07-10	Female	East Mark	2024-04-11	
120	1991-11-03	Male	East Charlotte	2024-03-02	
129	1953-09-28	Female	Lake Amystad	2024-01-05	
147	1980-06-26	Male	West Jennifer	2024-03-27	
175	1986-07-07	Female	Lake Codymouth	2024-02-13	
401	2004-06-22	Female	Port Charles	2024-02-27	
443	1957-04-24	Male	Gonzalezland	2024-01-08	

	Membership End Date	Subscription Plan	Payment Information	Renewal Status	\
6	2025-04-08	Monthly	Amex	Auto-renew	
33	2025-03-30	Annual	Mastercard	Auto-renew	
84	2025-02-02	Annual	Amex	Manual	
107	2025-04-11	Monthly	Amex	Manual	
120	2025-03-02	Annual	Visa	Manual	
129	2025-01-04	Annual	Mastercard	Manual	
147	2025-03-27	Monthly	Visa	Manual	
175	2025-02-12	Monthly	Mastercard	Auto-renew	
401	2025-02-26	Annual	Amex	Auto-renew	
443	2025-01-07	Annual	Amex	Manual	

	Usage Frequency	Purchase History	Favorite Genres	Devices Used	\
6	Frequent	Clothing	Sci-Fi	Tablet	
33	Regular	Electronics	Horror	Smartphone	
84	Occasional	Clothing	Horror	Smartphone	
107	Regular	Clothing	Action	Tablet	
120	Frequent	Books	Drama	Smartphone	
129	Occasional	Electronics	Sci-Fi	Tablet	
147	Regular	Books	Horror	Smart TV	
175	Frequent	Books	Horror	Smart TV	
401	Frequent	Electronics	Documentary	Smartphone	
443	Frequent	Books	Romance	Smart TV	

	Engagement Metrics	Feedback/Ratings	Customer Support Interactions
6	Medium	4.4	10
33	Low	4.5	10
84	Medium	3.0	1
107	Medium	3.2	7
120	High	4.2	2
129	Medium	4.4	3
147	Low	4.3	0
175	High	3.3	8

401	High	4.0	6
443	Medium	3.5	5

```
[11]: # Remove columns that I won't need
data = data.drop(columns = ['User ID', 'Name', 'Email Address', 'Username'])
```

```
[12]: # Create an age column based on years since the date of birth
current_date = datetime.now()
current_year = current_date.year
data['Age'] = current_year - data['Date of Birth'].dt.year
data.head()
```

```
[12]:
```

	Date of Birth	Gender	Location	Membership Start Date	\
0	1953-06-03	Male	Rebeccachester	2024-01-15	
1	1978-07-08	Male	Mcpersonview	2024-01-07	
2	1994-12-06	Female	Youngfort	2024-04-13	
3	1964-12-22	Female	Feliciashire	2024-01-24	
4	1961-06-04	Male	Port Deborah	2024-02-14	

	Membership End Date	Subscription Plan	Payment Information	Renewal Status	\
0	2025-01-14	Annual	Mastercard	Manual	
1	2025-01-06	Monthly	Visa	Manual	
2	2025-04-13	Monthly	Mastercard	Manual	
3	2025-01-23	Monthly	Amex	Auto-renew	
4	2025-02-13	Annual	Visa	Auto-renew	

	Usage Frequency	Purchase History	Favorite Genres	Devices Used	\
0	Regular	Electronics	Documentary	Smart TV	
1	Regular	Electronics	Horror	Smartphone	
2	Regular	Books	Comedy	Smart TV	
3	Regular	Electronics	Documentary	Smart TV	
4	Frequent	Clothing	Drama	Smart TV	

	Engagement Metrics	Feedback/Ratings	Customer Support Interactions	Age
0	Medium	3.6	3	71
1	Medium	3.8	7	46
2	Low	3.3	8	30
3	High	3.3	7	60
4	Low	4.3	1	63

```
[13]: # Check month distribution for memberships started
data['Month'] = data['Membership Start Date'].dt.month
data['Month_Text'] = data['Membership Start Date'].dt.month_name().str.
↳ slice(stop=3)

# Group by `Month` and `Month_Text`, sum it, and sort. Assign result to new
↳ DataFrame
```

```
data_by_month = data.groupby(['Month', 'Month_Text']).size().
    ↪reset_index(name='Count').sort_values('Month').head(12)
data_by_month

# April is lower than the rest as the data only goes through 4-13-2024
```

```
[13]:
```

	Month	Month_Text	Count
0	1	Jan	773
1	2	Feb	651
2	3	Mar	744
3	4	Apr	332

3 Exploratory Data Analysis (EDA)

```
[15]: # Create subplots for two plots
fig, axes = plt.subplots(nrows=1, ncols=2, figsize=(12,6))

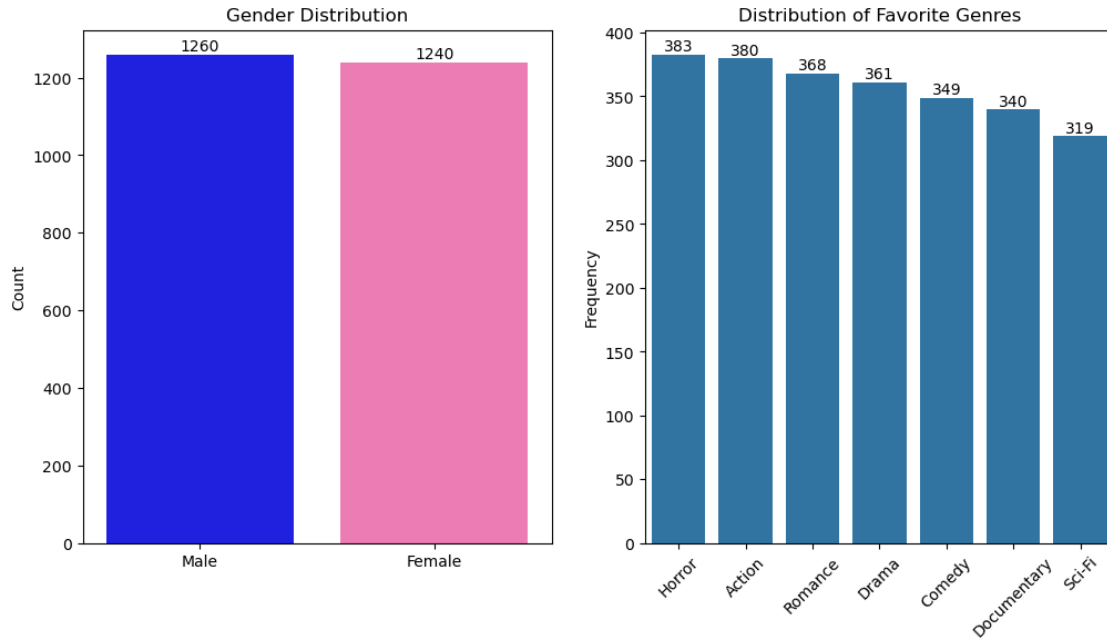
# Plot 1: Bar plot of Gender Distribution
gender_counts = data['Gender'].value_counts()
sns.barplot(x=gender_counts.index, y=gender_counts.values, ax=axes[0],
    ↪palette=['blue', 'hotpink'])
axes[0].set_title('Gender Distribution')
axes[0].set_xlabel('')
axes[0].set_ylabel('Count')

# Add counts above the bars
for index, value in enumerate(gender_counts.values):
    axes[0].text(index, value + 0.1, str(value), ha='center', va='bottom')

# Plot 2: Bar plot of Favorite Genres
genre_counts = data['Favorite Genres'].value_counts()
sns.barplot(x=genre_counts.index, y=genre_counts.values, ax=axes[1])
axes[1].set_title('Distribution of Favorite Genres')
axes[1].set_xlabel('')
axes[1].set_ylabel('Frequency')
axes[1].tick_params(axis='x', rotation=45) # Rotate x-axis labels for better
    ↪readability

# Add counts above the bars
for index, value in enumerate(genre_counts.values):
    axes[1].text(index, value + 0.1, str(value), ha='center', va='bottom')

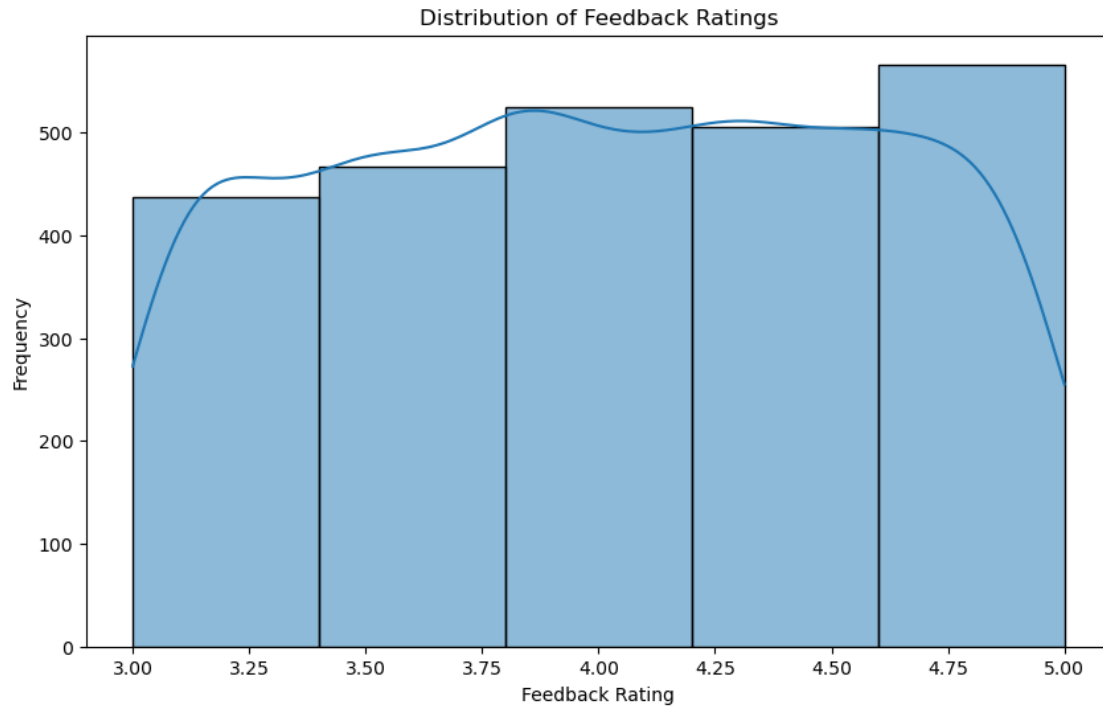
plt.show()
```



```
[16]: # Plot: Histogram of Feedback Ratings
plt.figure(figsize=(10,6))

plot = sns.histplot(data=data, x='Feedback/Ratings', bins=5, kde=True)
plot.set_title('Distribution of Feedback Ratings')
plot.set_xlabel('Feedback Rating')
plot.set_ylabel('Frequency')


plt.show()
```

```
[17]: # Group data by gender and favorite genres, and calculate counts
genre_summary = data.groupby(['Gender', 'Favorite Genres']).size().
    ↪reset_index(name='Count')
#print(genre_summary)

# Pivot the DataFrame to have favorite genres as columns
genre_summary_pivot = genre_summary.pivot(index='Favorite Genres',
    ↪columns='Gender', values='Count').fillna(0)

# Plot
genre_summary_pivot.plot(kind='bar', figsize=(10, 6), color=['hotpink', 'blue'])
plt.title('Genre Preference by Gender')
plt.xlabel('')
plt.ylabel('')
plt.xticks(rotation=45, ha='right')
plt.legend(title='Gender')
plt.tight_layout()
plt.show()
```



```
[18]: # Check the options and amounts of usage frequency
data['Usage Frequency'].value_counts()
```

```
[18]: Usage Frequency
      Frequent      851
      Regular      827
      Occasional    822
      Name: count, dtype: int64
```

4 Model Development

```
[20]: # Create a copy of the dataset to manipulate for the model
data_model = data.copy()

# Drop more features for easier reading and usage, now that the initial EDA is
↳ complete
data_model = data_model.drop(columns = ['Gender',
                                         'Location',
                                         'Membership Start Date',
                                         'Membership End Date',
                                         'Date of Birth',
                                         'Month_Text',
                                         'Month',
```

```
'Feedback/Ratings',  
'Payment Information',  
'Age']])
```

```
[21]: # Define a mapping dictionary for usage freq.  
usage_frequency_mapping = {  
    'Occasional': 1,  
    'Regular': 2,  
    'Frequent': 3  
}  
  
# Define a mapping dictionary for engagement metrics  
engagement_metrics_mapping = {  
    'Low': 1,  
    'Medium': 2,  
    'High': 3  
}  
  
# Apply the mapping to the "Usage Frequency" and "Engagement Metrics"  
data_model['Usage Frequency'] = data_model['Usage Frequency'].  
    ↪map(usage_frequency_mapping)  
  
data_model['Engagement Metrics'] = data_model['Engagement Metrics'].  
    ↪map(engagement_metrics_mapping)
```

```
[22]: # Define the columns to convert to dummy variables  
columns_to_convert = ['Renewal Status',  
    'Devices Used',  
    'Purchase History',  
    'Favorite Genres',  
    'Customer Support Interactions']  
  
# Create dummy variables for the specified columns  
dummy_variables = pd.get_dummies(data_model[columns_to_convert],  
    ↪drop_first=True, prefix='', prefix_sep='', dtype=int)  
  
# Concatenate the dummy variables with the original dataset  
data_model = pd.concat([data_model, dummy_variables], axis=1)  
  
# Drop the original categorical columns  
data_model.drop(columns=columns_to_convert, inplace=True)
```

```
[23]: # Check back on the data set  
data_model.head()
```

[23]:	Subscription Plan	Usage Frequency	Engagement Metrics	Manual	Smartphone	\
0	Annual	2	2	1	0	
1	Monthly	2	2	1	1	
2	Monthly	2	1	1	0	
3	Monthly	2	3	0	0	
4	Annual	3	1	0	0	

	Tablet	Clothing	Electronics	Comedy	Documentary	Drama	Horror	Romance	\
0	0	0	1	0	1	0	0	0	
1	0	0	1	0	0	0	1	0	
2	0	0	0	1	0	0	0	0	
3	0	0	1	0	1	0	0	0	
4	0	1	0	0	0	1	0	0	

	Sci-Fi
0	0
1	0
2	0
3	0
4	0

```
[24]: # Split the data into features (X) and target variable (y)
X = data_model.drop(columns=['Subscription Plan']) # Features
y = data_model['Subscription Plan'] # Target variable

# Split the data into training and testing sets (80% train, 20% test)
X_train, X_test, y_train, y_test = train_test_split(X, y, train_size = .
↪8, test_size=0.2, random_state=42)

# Define a list of models
models = [
    LogisticRegression(max_iter=1000, random_state=42),
    DecisionTreeClassifier(random_state=42),
    RandomForestClassifier(random_state=42)
]

# Iterate over each model
for model in models:

    # Train the model on the training data
    model.fit(X_train, y_train)

    # Make predictions on the testing data
    y_pred = model.predict(X_test)
```

```

# Evaluate the model's performance
accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred, pos_label='Annual')
recall = recall_score(y_test, y_pred, pos_label='Annual')
f1 = f1_score(y_test, y_pred, pos_label = "Annual")

# Print the model's performance
print(f"Model: {model.__class__.__name__}, Accuracy: {accuracy:.5f},  

↳ Precision: {precision:.5f}, Recall: {recall:.5f}, F1: {f1:.5f}")

```

Model: LogisticRegression, Accuracy: 0.52200, Precision: 0.51515, Recall: 0.68273, F1: 0.58722

Model: DecisionTreeClassifier, Accuracy: 0.51800, Precision: 0.51418, Recall: 0.58233, F1: 0.54614

Model: RandomForestClassifier, Accuracy: 0.53800, Precision: 0.53689, Recall: 0.52610, F1: 0.53144

```

[25]: # Retest the data to determine if a different split improves accuracy

# Split the data into training and testing sets (70% train, 30% test)
X_train, X_test, y_train, y_test = train_test_split(X, y, train_size=.7,  

↳ test_size=0.3, random_state=42)

models = [
    LogisticRegression(max_iter=1000, random_state=42),
    DecisionTreeClassifier(random_state=42),
    RandomForestClassifier(random_state=42)
]

# Iterate over each model
for model in models:

    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)
    accuracy = accuracy_score(y_test, y_pred)
    precision = precision_score(y_test, y_pred, pos_label='Annual')
    recall = recall_score(y_test, y_pred, pos_label='Annual')
    f1 = f1_score(y_test, y_pred, pos_label = "Annual")

    # Print the model's performance
    print(f"Model: {model.__class__.__name__}, Accuracy: {accuracy:.5f},  

↳ Precision: {precision:.5f}, Recall: {recall:.5f}, F1: {f1:.5f}")

```

Model: LogisticRegression, Accuracy: 0.49200, Precision: 0.51382, Recall: 0.56743, F1: 0.53930

Model: DecisionTreeClassifier, Accuracy: 0.54667, Precision: 0.56265, Recall: 0.60560, F1: 0.58333


Model: RandomForestClassifier, Accuracy: 0.52800, Precision: 0.55256, Recall:

0.52163, F1: 0.53665

```
[26]: # The Decision tree has the highest accuracy, precision, and recall.
      ↳ Outperforming the other models and performing better in the 70-30 split
      ↳ compared to the 80-20 split, most likely due to more data learn more
      ↳ effective.
```


```
[27]: # Calculate correlation matrix
corr_matrix = X.corr()

# Plot heatmap
plt.figure(figsize=(12, 10))
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm', fmt=".2f")
plt.title("Correlation Matrix")
plt.show()
```



```
[28]: # Get feature importances for random forest model
feature_importances = models[2].feature_importances_

# Plot feature importances
plt.figure(figsize=(10, 6))
plt.barh(X.columns, feature_importances)
plt.xlabel('Feature Importance')
plt.title('Feature Importance Scores')
plt.show()
```



4.0.1 Testing new hyperparameters utilizing GridSearch to see if a better random forest can be created

```
[30]: # Separate into labels and features
y = data_model['Subscription Plan']
X = data_model.drop('Subscription Plan', axis=1)

# train, val, test sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=.25,
    ↪random_state = 42)
X_tr, X_val, y_tr, y_val = train_test_split(X_train, y_train, test_size = 0.25,
    ↪random_state = 42)

# Set hyperparameters
cv_params = {'n_estimators' : [50,100],
             'max_depth' : [10,50],
```

```

        'min_samples_leaf' : [0.5,1],
        'min_samples_split' : [0.001, 0.01],
        'max_features' : ["sqrt"],
        'max_samples' : [.5,.9]}

# Split index
split_index = [0 if x in X_val.index else -1 for x in X_train.index]
custom_split = PredefinedSplit(split_index)

# Instantiate the model
rf = RandomForestClassifier(random_state=42)

# search over parameters
rf_val = GridSearchCV(rf, cv_params, cv=custom_split, refit='f1', n_jobs = -1,
    ↪ verbose = 1)

# Fit the model
rf_val.fit(X_train, y_train)

```

Fitting 1 folds for each of 32 candidates, totalling 32 fits

```

[30]: GridSearchCV(cv=PredefinedSplit(test_fold=array([-1, -1, ..., -1, -1])),
    estimator=RandomForestClassifier(random_state=42), n_jobs=-1,
    param_grid={'max_depth': [10, 50], 'max_features': ['sqrt'],
        'max_samples': [0.5, 0.9],
        'min_samples_leaf': [0.5, 1],
        'min_samples_split': [0.001, 0.01],
        'n_estimators': [50, 100]},
    refit='f1', verbose=1)

```

```

[31]: #obtain optimal paramaters
rf_val.best_params_

```

```

[31]: {'max_depth': 50,
    'max_features': 'sqrt',
    'max_samples': 0.9,
    'min_samples_leaf': 1,
    'min_samples_split': 0.001,
    'n_estimators': 100}

```

```

[32]: # Use optimal parameters on GridSearchCV.
rf_opt = RandomForestClassifier(n_estimators = 100, max_depth = 50,
    min_samples_leaf = 1, min_samples_split = 0.001,
    max_features="sqrt", max_samples = 0.9,
    ↪ random_state = 42)

```

```

[33]: # Fit the optimal model.

```



```
rf_opt.fit(X_train, y_train)
```

```
[33]: RandomForestClassifier(max_depth=50, max_samples=0.9, min_samples_split=0.001,
                             random_state=42)
```

```
[34]: # Predict on test set
y_pred = rf_opt.predict(X_test)

# Get scores
rf_pc_test = precision_score(y_test, y_pred, pos_label = "Annual")
rf_rc_test = recall_score(y_test, y_pred, pos_label = "Annual")
rf_ac_test = accuracy_score(y_test, y_pred)
rf_f1_test = f1_score(y_test, y_pred, pos_label = "Annual")

# Print results
table = pd.DataFrame({'Model': ["Decision Tree", "Tuned Random Forest"],
                       'F1': [0.58333, rf_f1_test],
                       'Recall': [0.60560, rf_rc_test],
                       'Precision': [0.56265, rf_pc_test],
                       'Accuracy': [0.54667, rf_ac_test]
                      })

table
```

```
[34]:
```

	Model	F1	Recall	Precision	Accuracy
0	Decision Tree	0.583330	0.605600	0.562650	0.54667
1	Tuned Random Forest	0.547655	0.555215	0.540299	0.52160

4.0.2 We can see that the Decision Tree with 75/25 split has better scores all around

```
[36]: data_model_binary = pd.get_dummies(data_model, drop_first=True, dtype=int)
data_model_binary.head()
```

```
[36]:
```

	Usage Frequency	Engagement Metrics	Manual	Smartphone	Tablet	Clothing	\
0	2	2	1	0	0	0	
1	2	2	1	1	0	0	
2	2	1	1	0	0	0	
3	2	3	0	0	0	0	
4	3	1	0	0	0	1	

	Electronics	Comedy	Documentary	Drama	Horror	Romance	Sci-Fi	\
0	1	0	1	0	0	0	0	
1	1	0	0	0	1	0	0	
2	0	1	0	0	0	0	0	
3	1	0	1	0	0	0	0	
4	0	0	0	1	0	0	0	

	Subscription Plan_Monthly
0	0
1	1
2	1
3	1
4	0

```
[37]: # XGBoost

# Separate into labels and features
y = data_model_binary['Subscription Plan_Monthly']
X = data_model_binary.drop('Subscription Plan_Monthly', axis=1)

# train, val, test sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=.25,
    random_state = 42)
X_tr, X_val, y_tr, y_val = train_test_split(X_train, y_train, test_size = 0.25,
    random_state = 42)

# Set hyperparameters
cv_params = {'learning_rate': [0.1],
             'max_depth': [8],
             'min_child_weight': [2],
             'n_estimators': [500]
            }
xgb = XGBClassifier(objective='binary:logistic', random_state=42)
```

```
[38]: scoring = {
    'accuracy': make_scorer(accuracy_score),
    'precision': make_scorer(precision_score),
    'recall': make_scorer(recall_score),
    'f1': make_scorer(f1_score),
}

xgb_cv = GridSearchCV(xgb, cv_params, scoring=scoring, cv=4, refit='f1')
```

```
[39]: %%time
# fit the GridSearch model to training data

xgb_cv.fit(X_train, y_train)
```

CPU times: total: 15.6 s
Wall time: 1.13 s

```
[39]: GridSearchCV(cv=4,
                  estimator=XGBClassifier(base_score=None, booster=None,
                                          callbacks=None, colsample_bylevel=None,
```

```

        colsample_bynode=None,
        colsample_bytree=None, device=None,
        early_stopping_rounds=None,
        enable_categorical=False, eval_metric=None,
        feature_types=None, gamma=None,
        grow_policy=None, importance_type=None,
        interaction_constraints=None,
        learning_rate=None,...
    param_grid={'learning_rate': [0.1], 'max_depth': [8],
                'min_child_weight': [2], 'n_estimators': [500]},
    refit='f1',
    scoring={'accuracy': make_scorer(accuracy_score,
response_method='predict'),
            'f1': make_scorer(f1_score, response_method='predict'),
            'precision': make_scorer(precision_score,
response_method='predict'),
            'recall': make_scorer(recall_score,
response_method='predict')}})

```

```

[40]: # Examine best score and paramaters
print(xgb_cv.best_score_)

print(xgb_cv.best_params_)

```

```

0.4903714701922771
{'learning_rate': 0.1, 'max_depth': 8, 'min_child_weight': 2, 'n_estimators':
500}

```

```

[41]: # Predict on test set
y_pred = xgb_cv.predict(X_test)

# Get scores
xgb_pc_test = precision_score(y_test, y_pred)
xgb_rc_test = recall_score(y_test, y_pred)
xgb_ac_test = accuracy_score(y_test, y_pred)
xgb_f1_test = f1_score(y_test, y_pred)

# Print results
table = pd.DataFrame({'Model': ["Decision Tree", "Tuned Random Forest", "XGBoost"],
                        'F1': [0.58333, rf_f1_test, xgb_pc_test],
                        'Recall': [0.60560, rf_rc_test, xgb_rc_test],
                        'Precision': [0.56265, rf_pc_test, xgb_ac_test],
                        'Accuracy': [0.54667, rf_ac_test, xgb_f1_test]
                      })

table

```

[41]:	Model	F1	Recall	Precision	Accuracy
0	Decision Tree	0.583330	0.605600	0.562650	0.546670
1	Tuned Random Forest	0.547655	0.555215	0.540299	0.521600
2	XGBoost	0.483871	0.501672	0.505600	0.492611

5 Model Selection

5.1 Results and Evaluation

Tree-based Machine Learning and Gradient Boosting

After conducting feature engineering, the decision tree model achieved f1-score of 58.3%, recall of 60.6%, precision of 56.3%, and accuracy of 54.7%, on the test set.

5.2 Conclusion, Recommendations, Next Steps

The models and the feature importances extracted from the models confirm that employees at the company are overworked.

To retain customers, the following recommendations could be presented to the stakeholders:

- **Personalized Subscription Plans:** Utilize the insights from the predictive models to offer personalized subscription plans tailored to individual user preferences, usage frequency, and engagement metrics. This can enhance user satisfaction and retention by providing plans that better suit their needs.
- **Promotional Campaigns:** Target promotional campaigns towards users who are more likely to churn based on the predictive models. Offer incentives or discounts to encourage them to renew their subscriptions or upgrade to higher-tier plans.
- **Content Recommendations:** Leverage the favorite genres and purchase history of users to provide targeted content recommendations. This can increase user engagement and satisfaction by delivering content that aligns with their interests.
- **Customer Support Enhancement:** Identify users who have had frequent interactions with customer support and proactively address their concerns or issues. Improving the customer support experience can lead to higher retention rates and positive feedback.
- **Optimized Renewal Strategies:** Implement strategies to optimize subscription renewal processes, such as providing seamless renewal options, offering incentives for automatic renewal, or sending timely reminders before subscription expiration dates.
- **Continuous Monitoring and Feedback:** Regularly monitor user feedback and engagement metrics to identify evolving preferences and trends. Incorporate user feedback into product development and subscription offerings to ensure continued relevance and satisfaction.
- **Further Model Refinement:** Continuously refine and update the predictive models based on new data and feedback. This can improve the accuracy and effectiveness of the models over time, leading to better decision-making and outcomes.
- **More data:** This model could be significantly improved with more datapoints.